




(a) Reward accuracy on masked spans




(b) Reward margin on masked spans



(c) FiRe-MPO loss decomposition (Qwen3-VL-4B)



(d) Gradient norm



— DPO — FiRe-MPO — HuatuoGPT-Vision-7B - - Qwen3-VL-4B